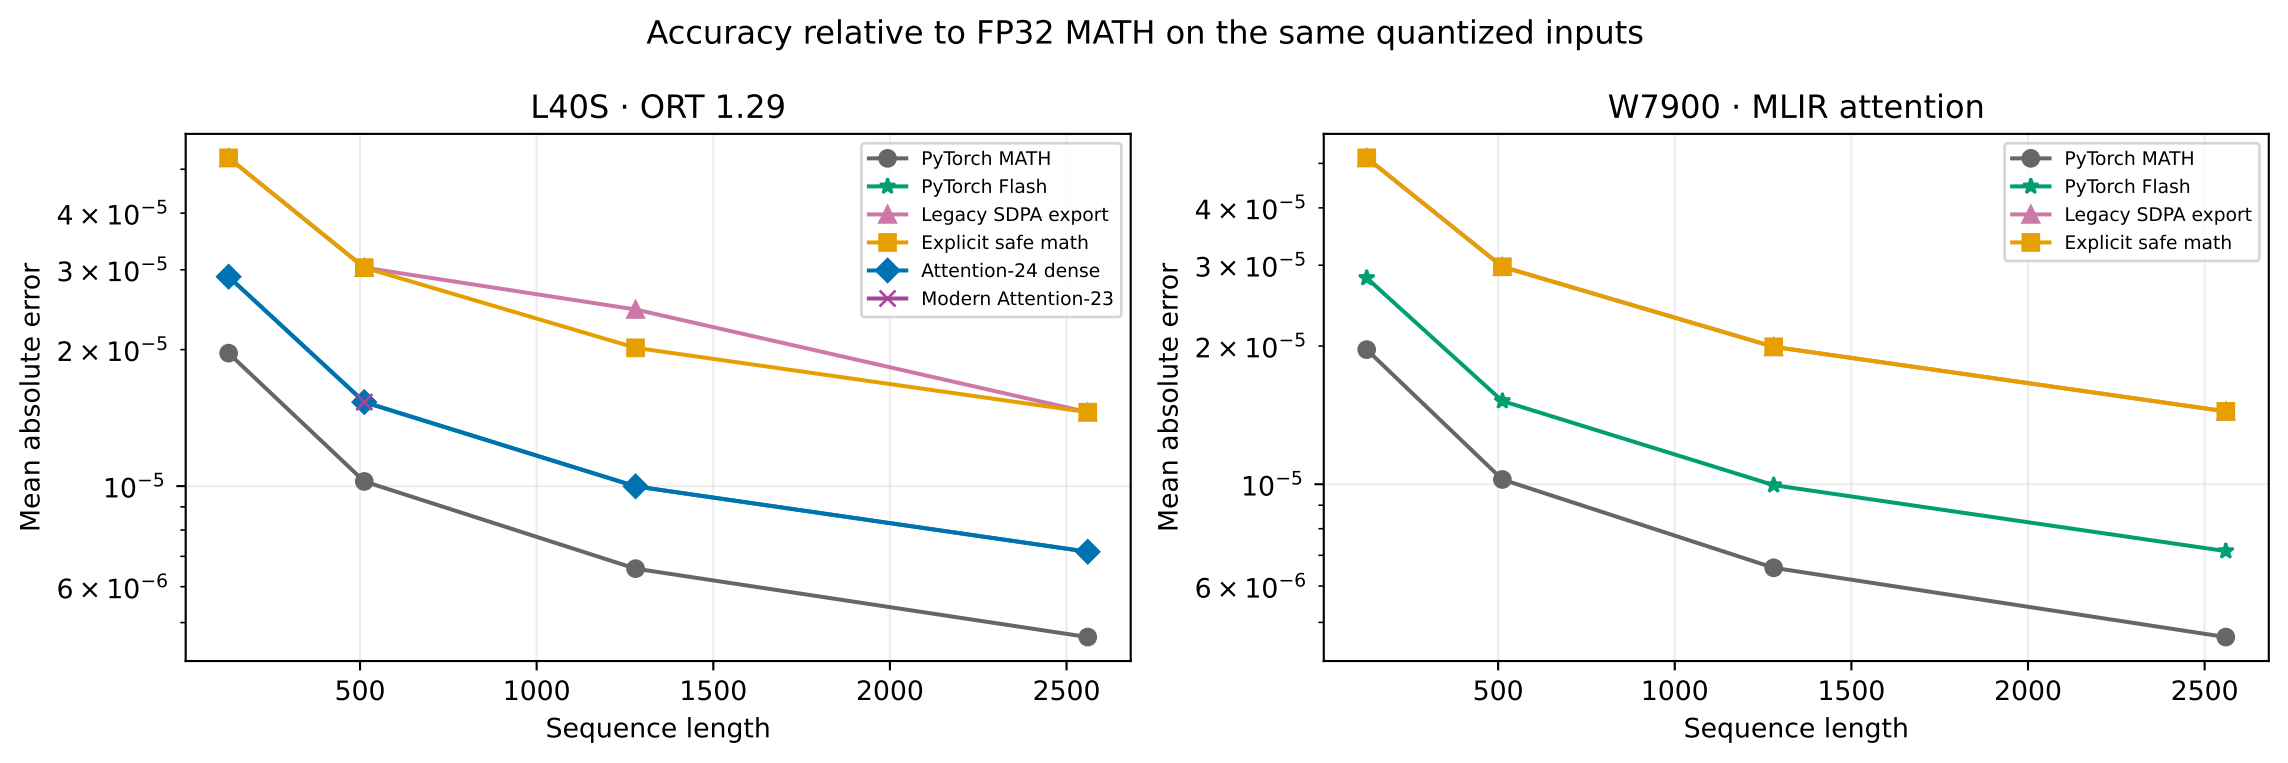
Accuracy relative to FP32 MATH on the same quantized inputs
L40S · ORT 1.29
W7900 · MLIR attention
Mean absolute error
Sequence length
PyTorch MATH
PyTorch Flash
Legacy SDPA export
Explicit safe math
Attention-24 dense
Modern Attention-23
4 × 10⁻⁵
3 × 10⁻⁵
2 × 10⁻⁵
10⁻⁵
6 × 10⁻⁶
500
1000
1500
2000
2500