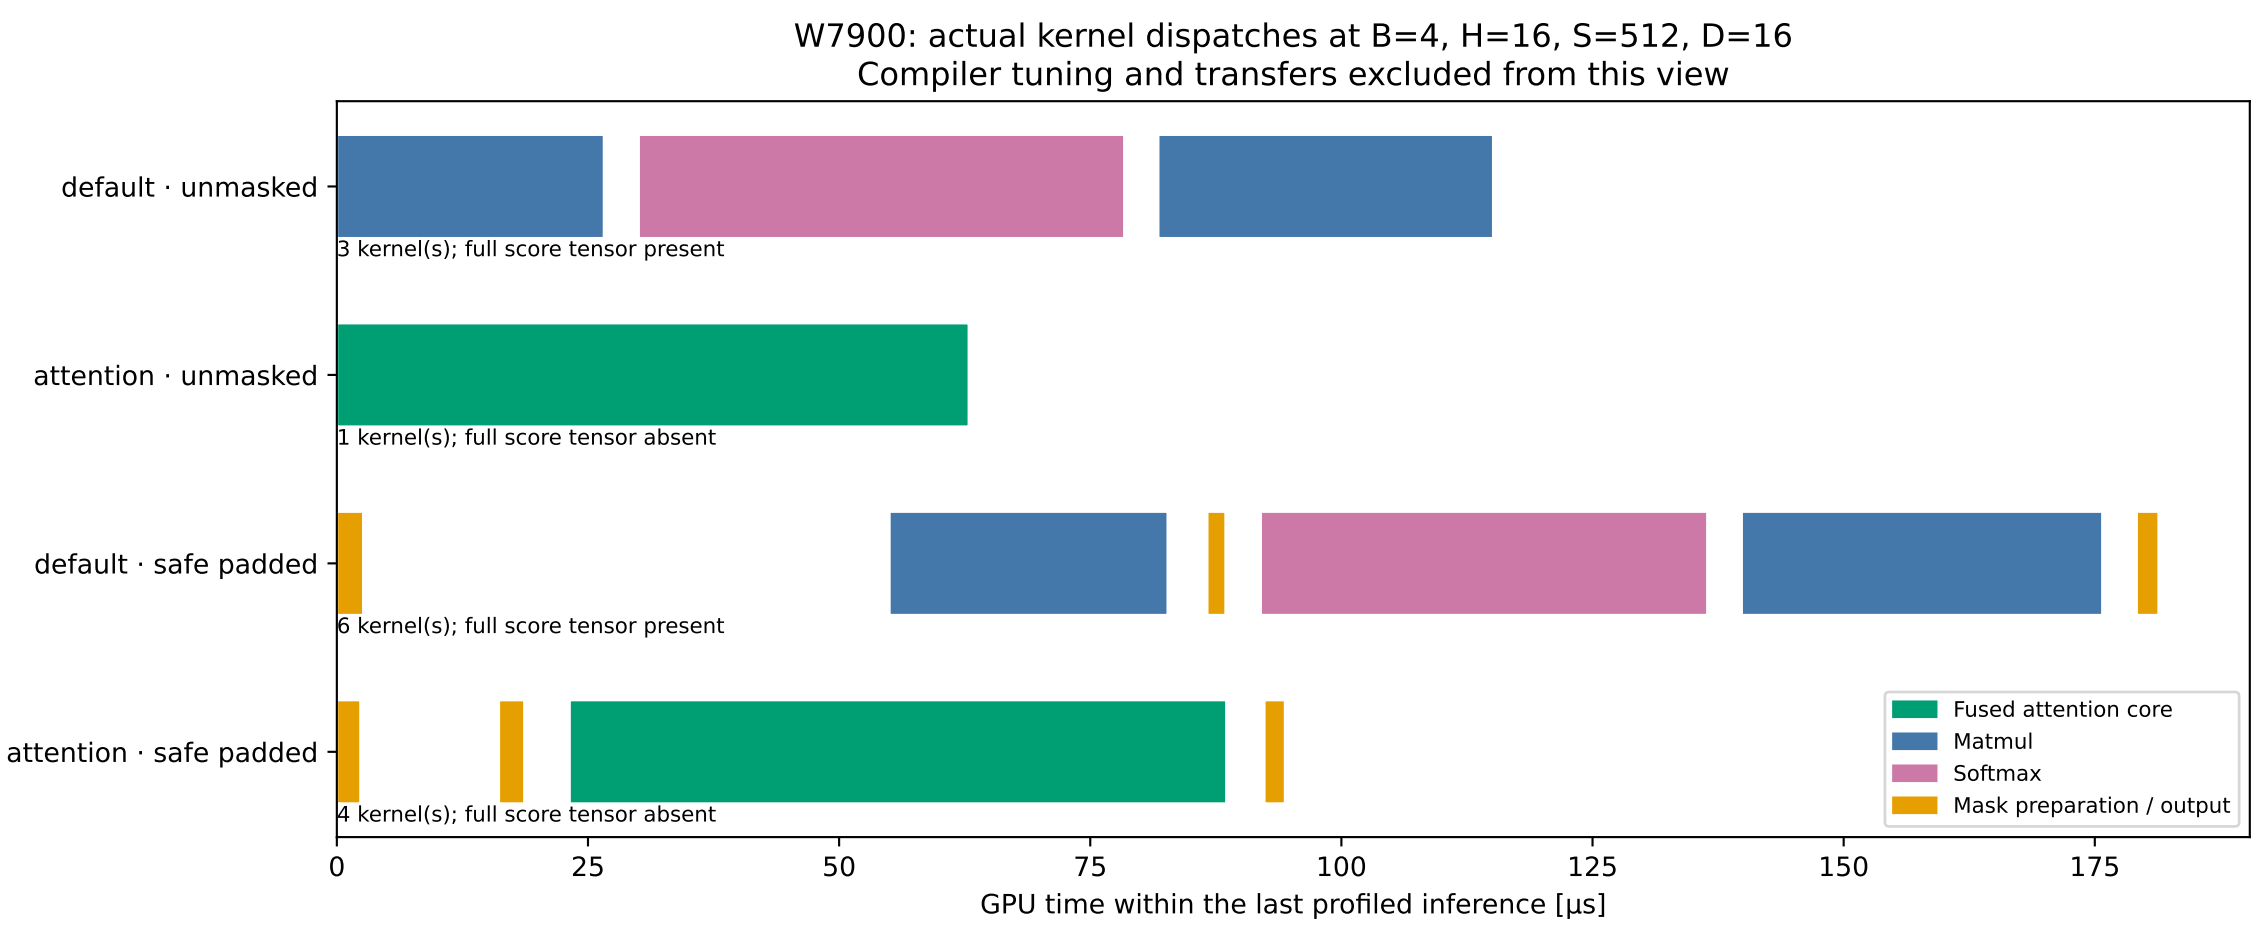
W7900: actual kernel dispatches at B=4, H=16, S=512, D=16
Compiler tuning and transfers excluded from this view
default · unmasked
3 kernel(s); full score tensor present
attention · unmasked
1 kernel(s); full score tensor absent
default · safe padded
6 kernel(s); full score tensor present
attention · safe padded
4 kernel(s); full score tensor absent
Fused attention core
Matmul
Softmax
Mask preparation / output
0
25
50
75
100
125
150
175
GPU time within the last profiled inference [µs]