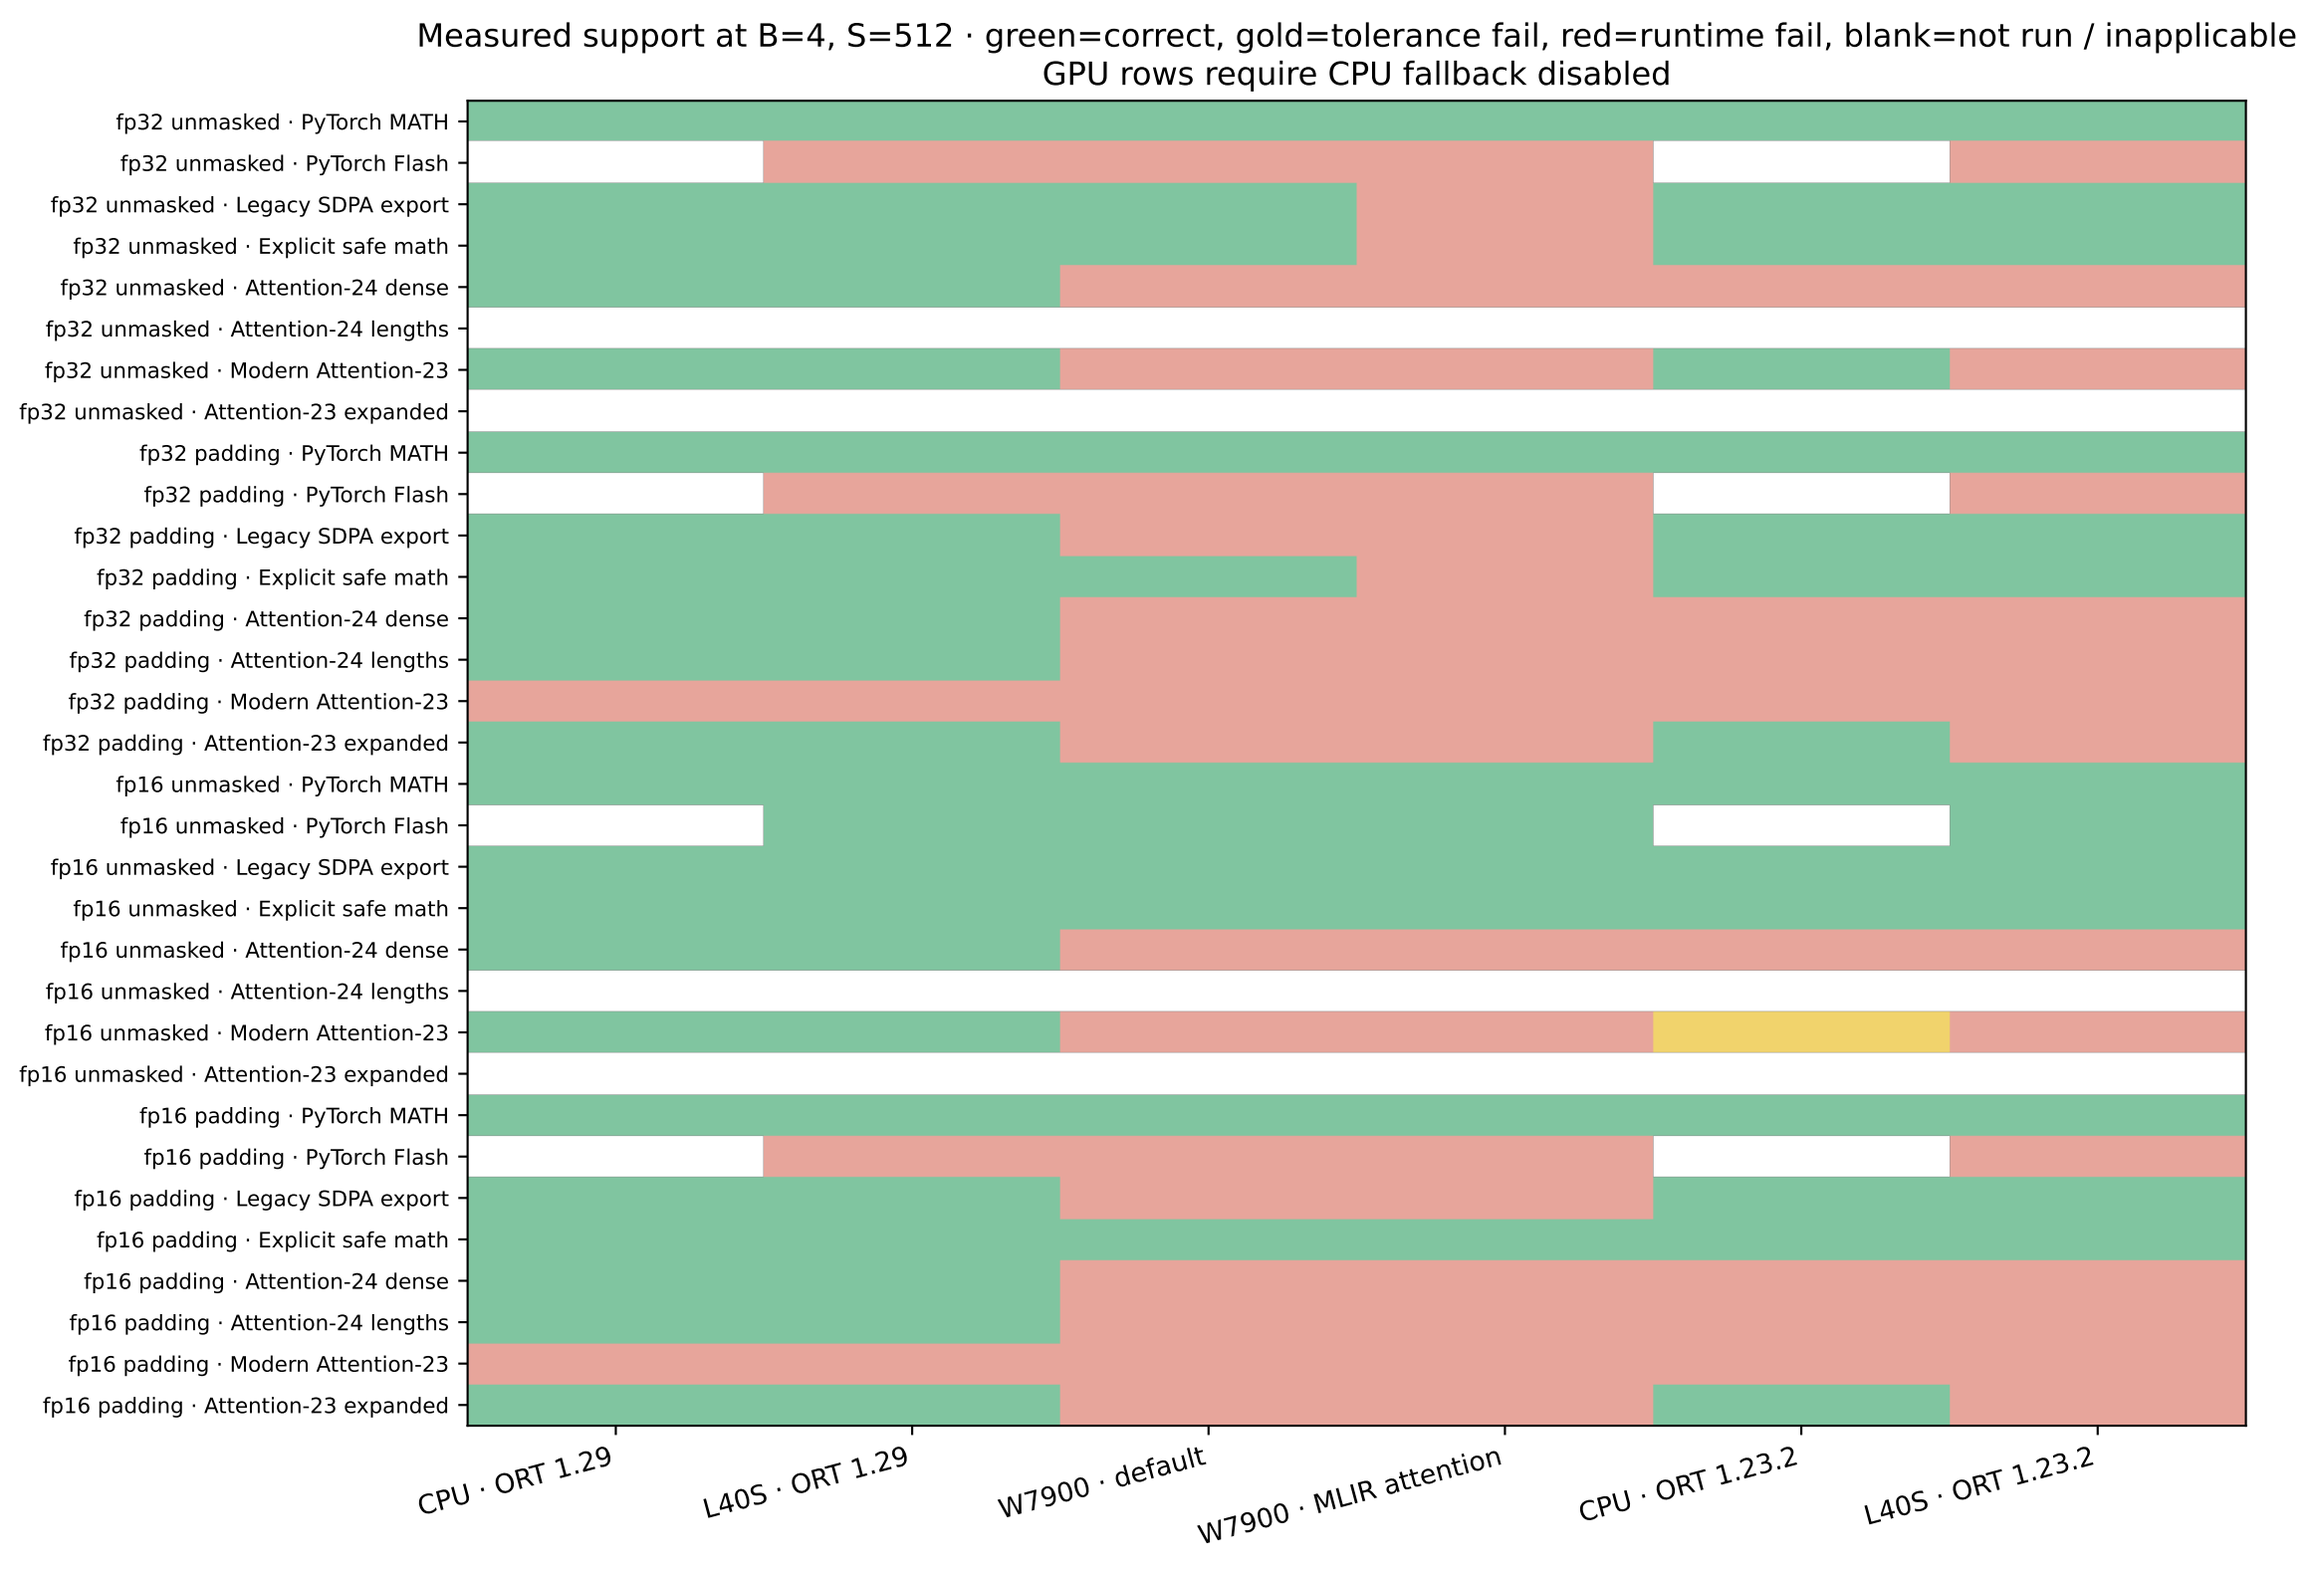
Measured support at B=4, S=512 · green=correct, gold=tolerance fail, red=runtime fail, blank=not run / inapplicable
GPU rows require CPU fallback disabled
fp32 unmasked · PyTorch MATH
fp32 unmasked · PyTorch Flash
fp32 unmasked · Legacy SDPA export
fp32 unmasked · Explicit safe math
fp32 unmasked · Attention-24 dense
fp32 unmasked · Attention-24 lengths
fp32 unmasked · Modern Attention-23
fp32 unmasked · Attention-23 expanded
fp32 padding · PyTorch MATH
fp32 padding · PyTorch Flash
fp32 padding · Legacy SDPA export
fp32 padding · Explicit safe math
fp32 padding · Attention-24 dense
fp32 padding · Attention-24 lengths
fp32 padding · Modern Attention-23
fp32 padding · Attention-23 expanded
fp16 unmasked · PyTorch MATH
fp16 unmasked · PyTorch Flash
fp16 unmasked · Legacy SDPA export
fp16 unmasked · Explicit safe math
fp16 unmasked · Attention-24 dense
fp16 unmasked · Attention-24 lengths
fp16 unmasked · Modern Attention-23
fp16 unmasked · Attention-23 expanded
fp16 padding · PyTorch MATH
fp16 padding · PyTorch Flash
fp16 padding · Legacy SDPA export
fp16 padding · Explicit safe math
fp16 padding · Attention-24 dense
fp16 padding · Attention-24 lengths
fp16 padding · Modern Attention-23
fp16 padding · Attention-23 expanded
CPU · ORT 1.29
L40S · ORT 1.29
W7900 · default
W7900 · MLIR attention
CPU · ORT 1.23.2
L40S · ORT 1.23.2