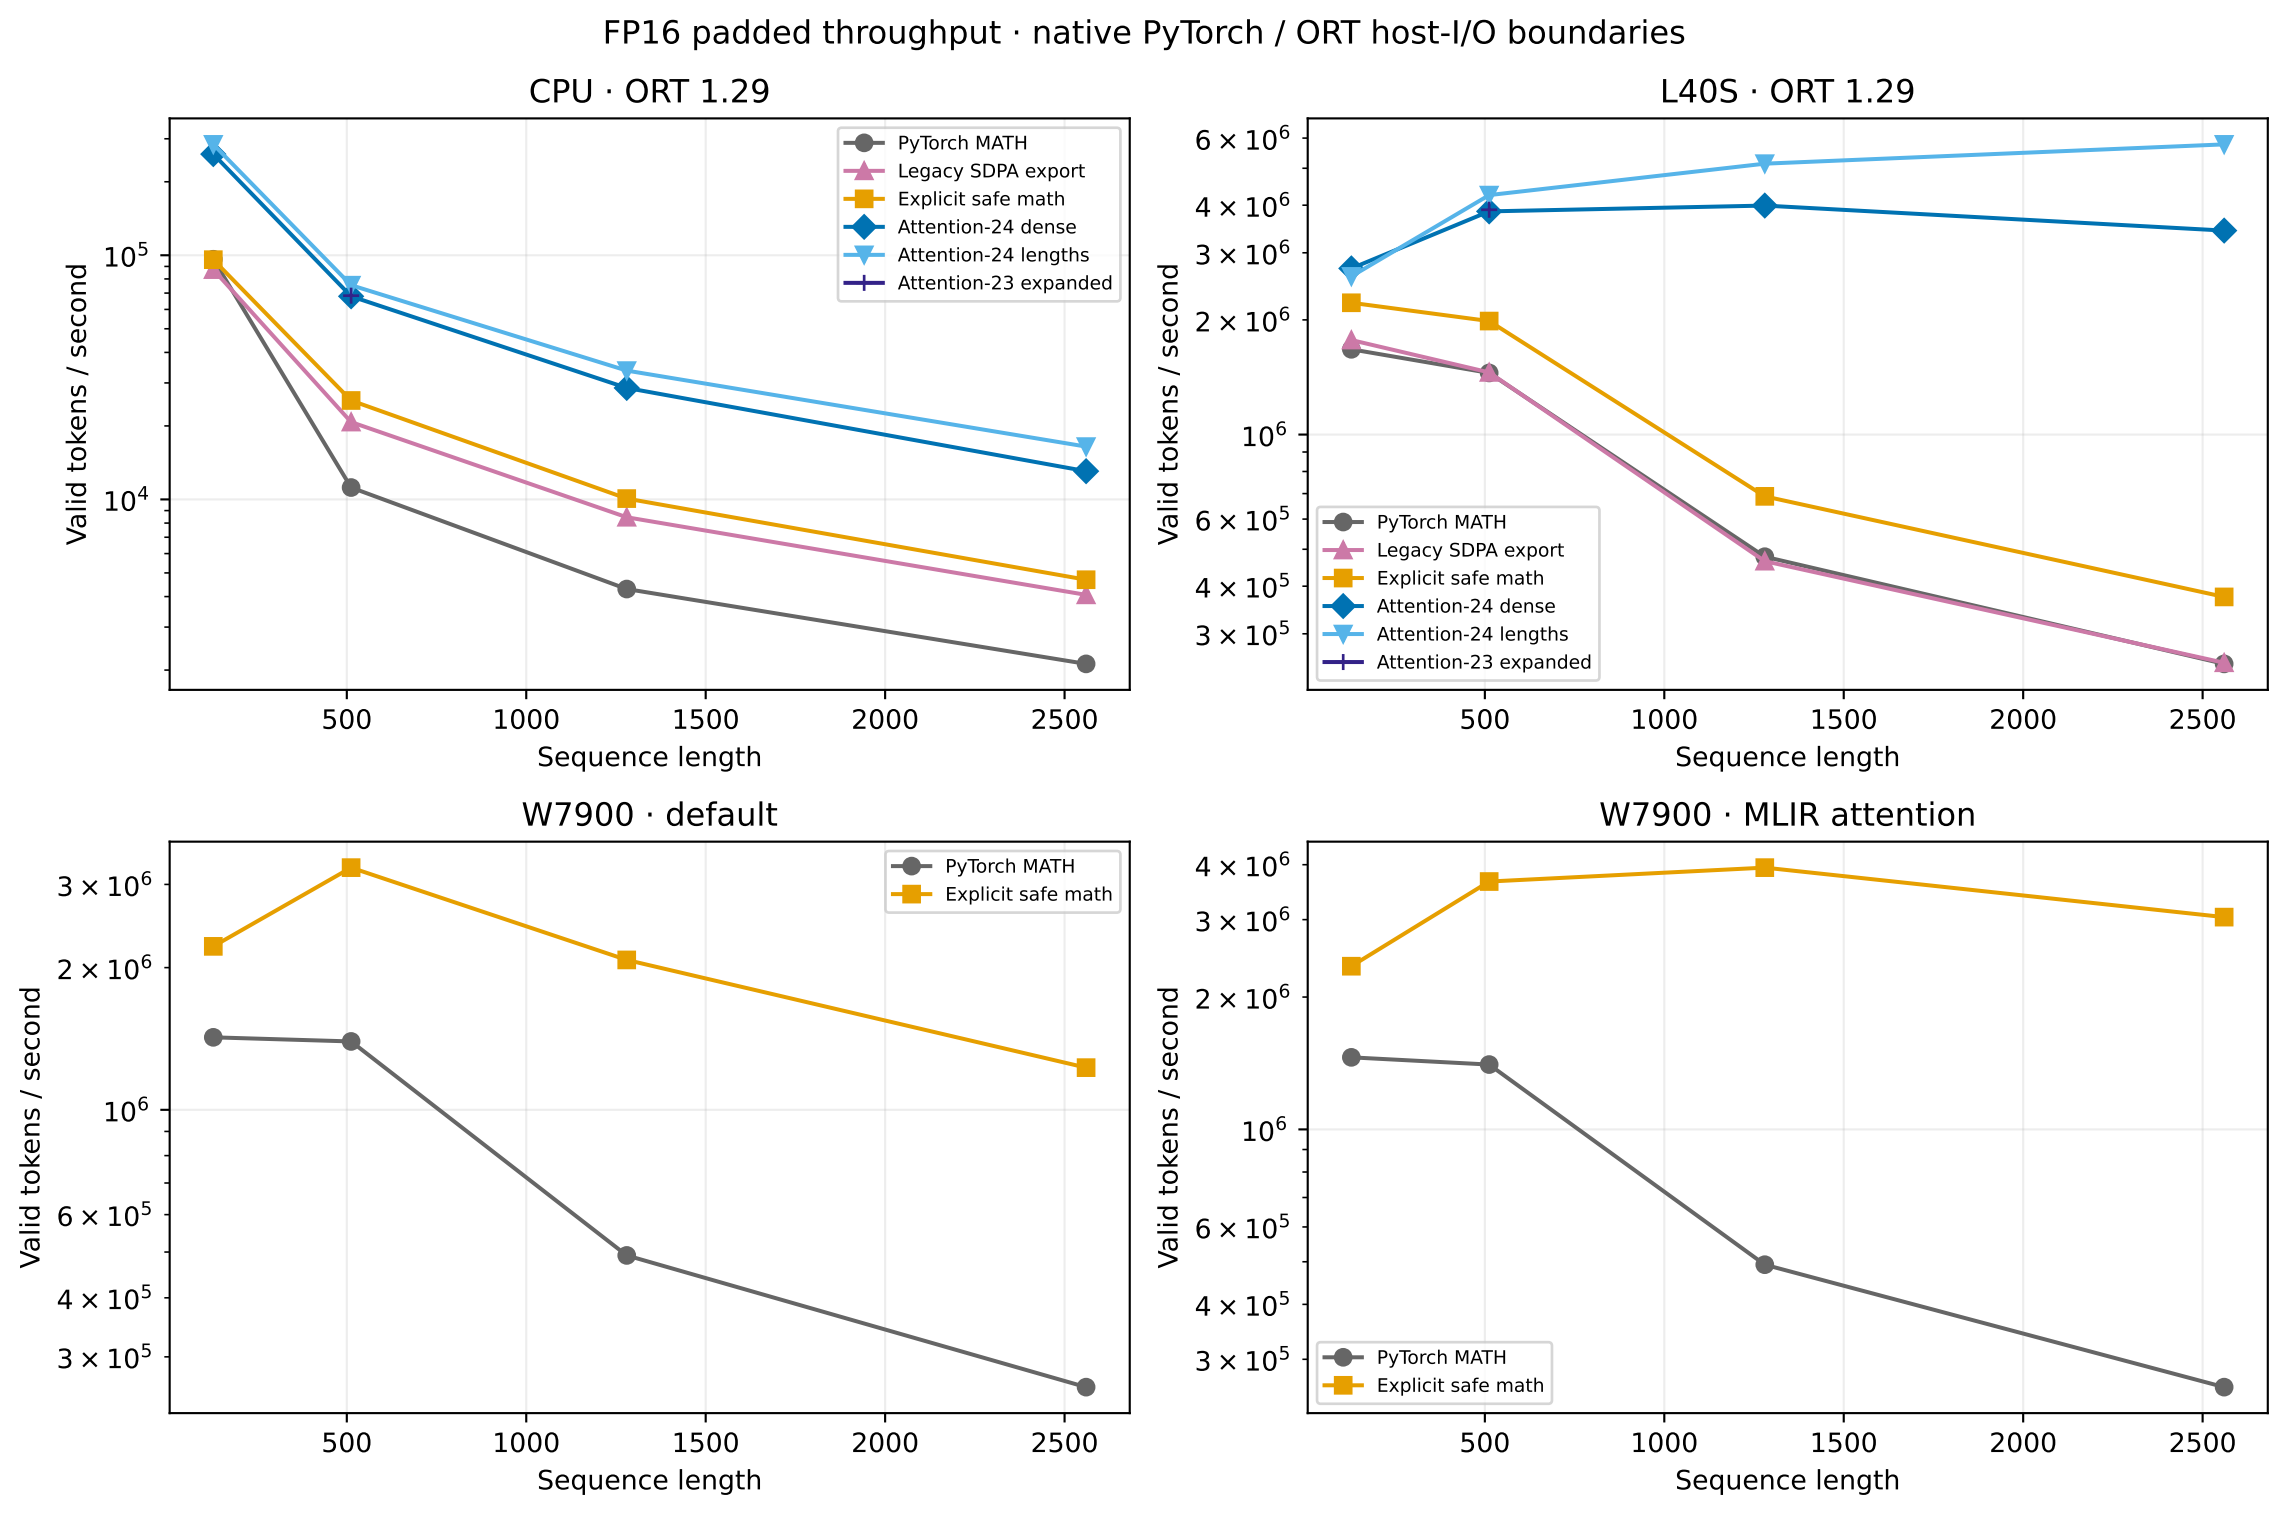
FP16 padded throughput · native PyTorch / ORT host-I/O boundaries
CPU · ORT 1.29
L40S · ORT 1.29
W7900 · default
W7900 · MLIR attention
Valid tokens / second
Sequence length
PyTorch MATH
Legacy SDPA export
Explicit safe math
Attention-24 dense
Attention-24 lengths
Attention-23 expanded